



EPL646 – Advanced Topics in Databases

Lecture 2

Storage I: Storage and Indexing

Chap. 8.1-8.5: Ramakr. & Gehrke

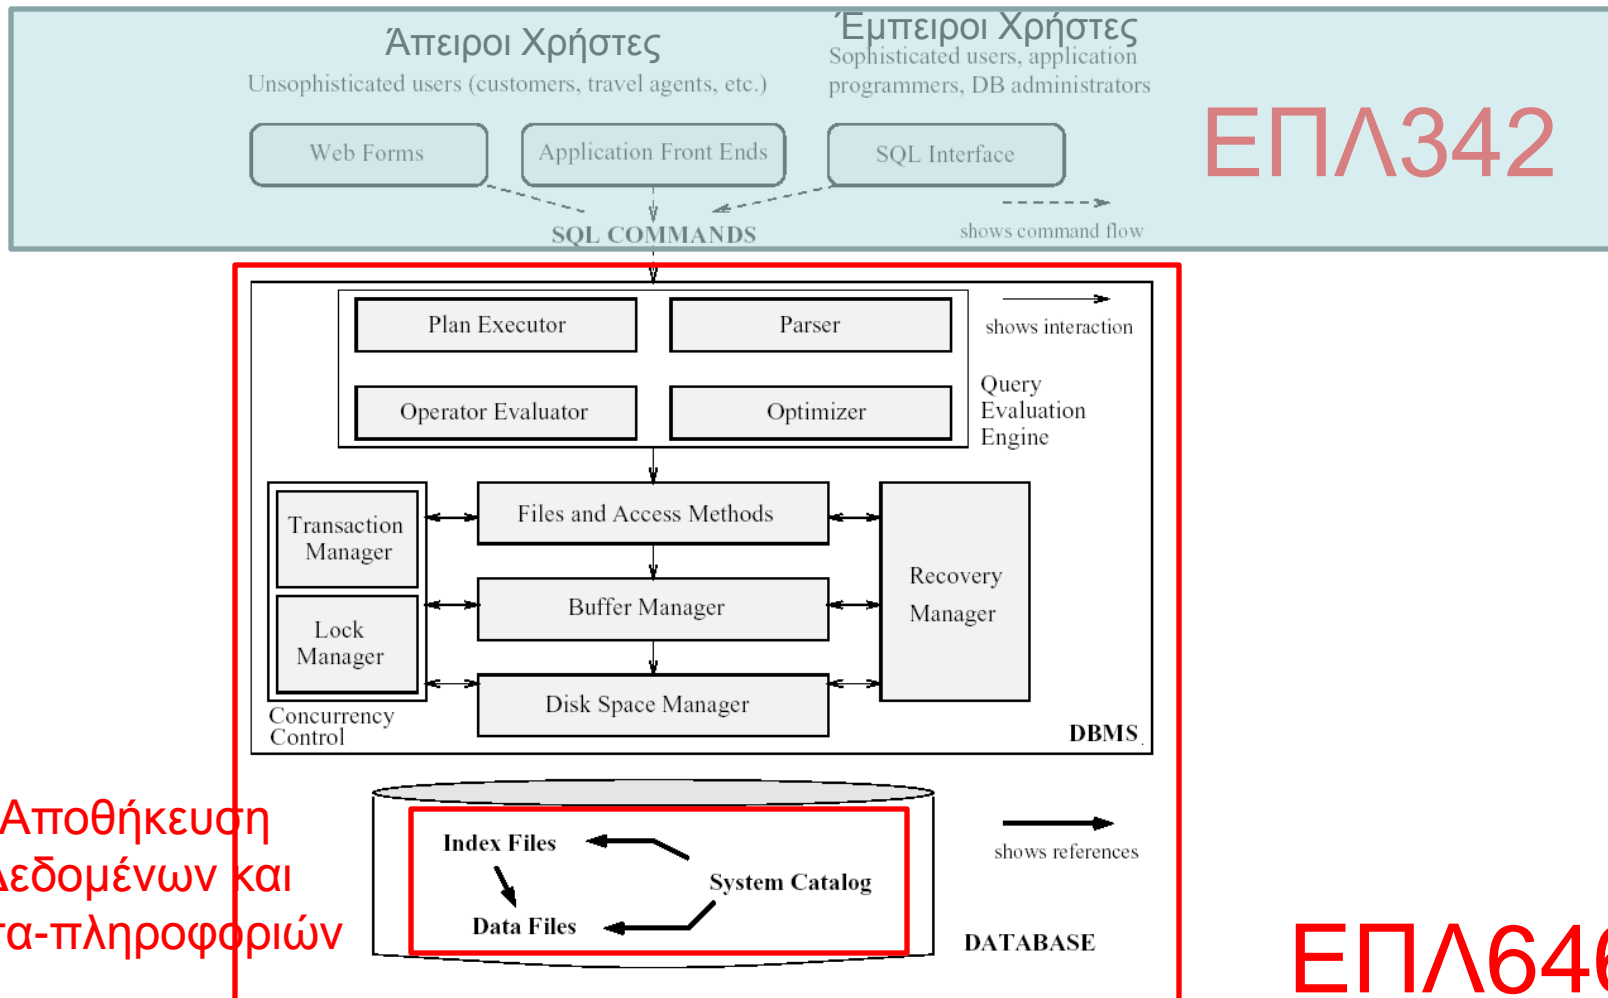
*** exclude 8.4.5-8.4.6**

Demetris Zeinalipour

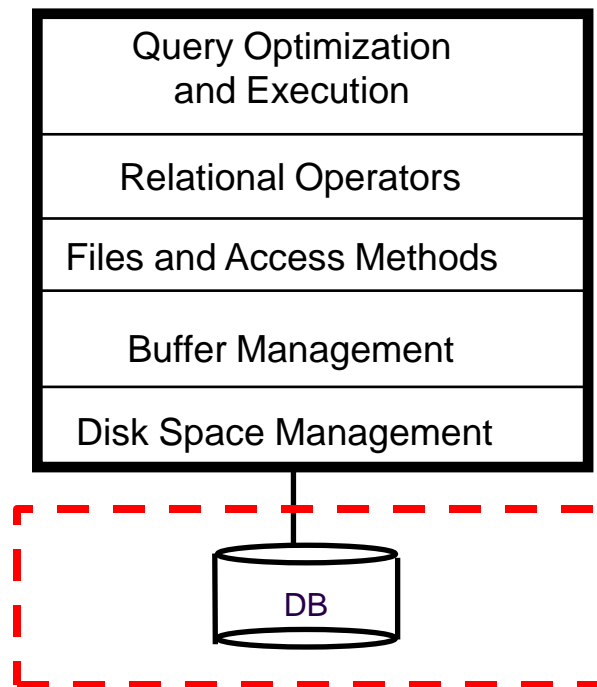
<http://www.cs.ucy.ac.cy/~dzeina/courses/epl646>

ΕΠΛ646: Ενότητα Α

Εσωτερική Λειτουργία ενός RDBMS



Context of next slides



Data on External Storage

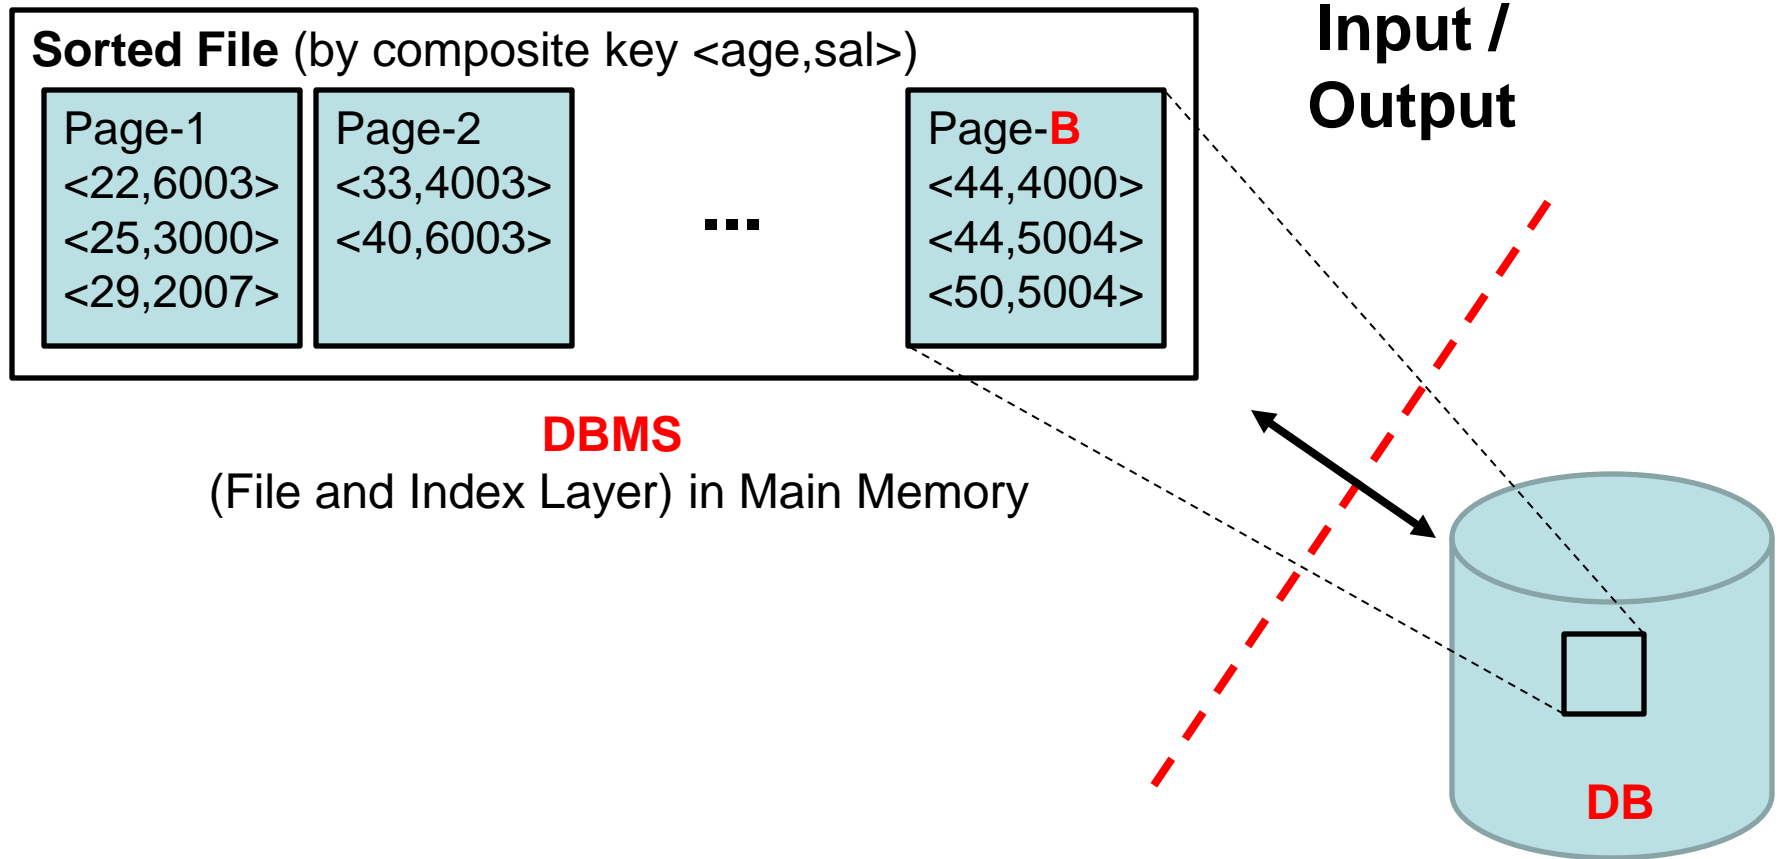
(Δεδομένα στη Δευτερεύουσα Μνήμη)



- A DBMS stores **vast amounts** of data and the data has to **persist** across program executions.
- Therefore, data is stored on **external storage** and **fetched** into **main memory** as needed for processing.
- The unit of information that is read and written to a disk is called **Page (Σελίδα)**, e.g., 4KB ή 8KB
- Higher layers of the DBMS view these pages as unified **Files (Αρχείο)** and can read/write **Records (Εγγραφές, Πλειάδες)** to these files.
 - Consider a data record (id:4B, name:28B) and a 4096B (4KB) page size. That would yield ~128 records / page (some bytes go to headers and other auxiliary structures).
- **What is the basic performance cost in a DBMS?**
 - **I/O (Input/Output): # pages read/write** for a given operation.
 - **Complexity** of algorithms in DBMSs is expressed in I/Os

Data on External Storage

(Δεδομένα στη Δευτερεύουσα Μνήμη)



Storage Mediums (Μέσα Αποθήκευσης)

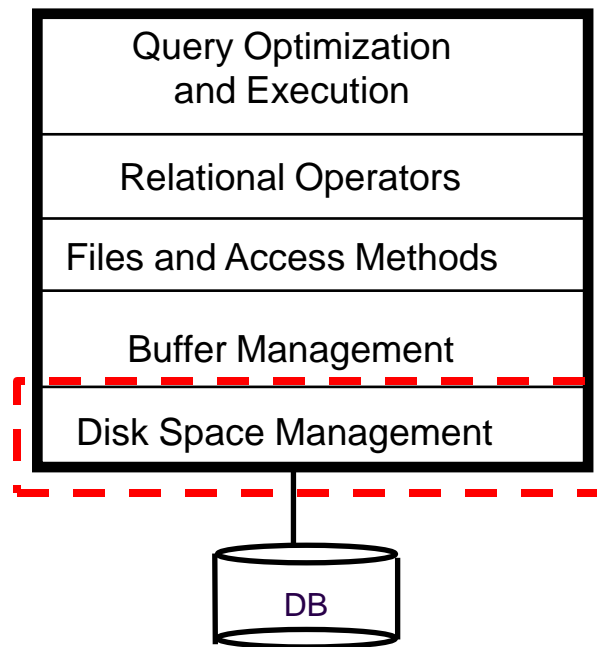


- Disks: Can retrieve a **random** page at a **fixed cost**
 - But reading several **consecutive pages** is much cheaper than reading them in **random order**
- Tapes: Can only read pages **in sequence**
 - **Cheaper** than disks; used for **archiving** (αρχειοθέτηση)
- Flash Memory (Solid State Disks): Reading data at the speed of main memory, writing is slower.
 - More **expensive** than disks; used for applications with read workloads that require **fast random** accesses.

Main focus
of DBMSs



Context of next slides



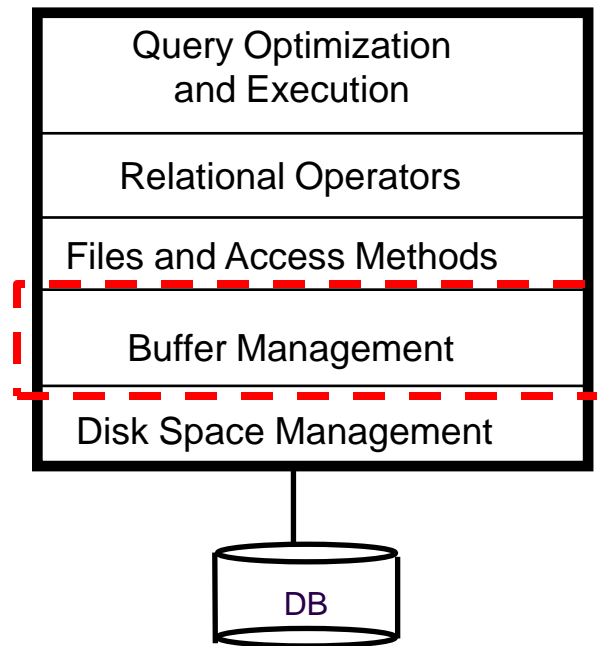
Disk Space Manager (DSM)

(Διαχειριστής Χώρου Δίσκου)



- **DSM:** Supports the concept of a **page** as a unit of data and provides commands to **allocate/deallocate**, **read/write** a page to external storage.
 - Size of Page == Size of Disk Block, in order to support read/write operations in one I/O operation.
 - The higher layers in the DB architecture (i.e., the **Buffer Manager**) interact directly with the DSM.
- Other Duties: **Keep track of Free Blocks.**
 - Initially a DB is stored on consecutive **disk blocks** (when it acts in its own partition) or inside a **file** (when it is stored inside an Operation System file).
 - Subsequent **deletions** might easily create “**holes**” in that sequence (either file or disk), thus the DSM needs to track the free pages.

Context of next slides



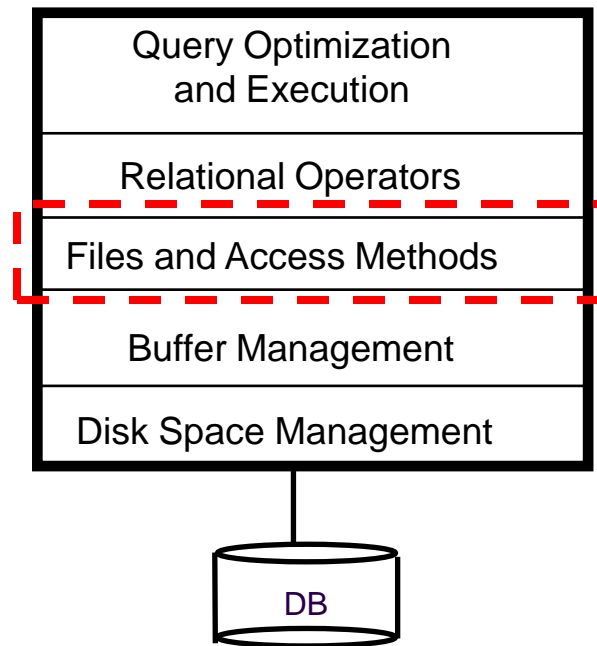
Buffer Manager (BM)

(Διαχειριστής Κρυφής Μνήμης)



- **BM:** Subsystem that is responsible for **loading** pages from **external storage** to the main memory **buffer pool**
 - File & Index layers make calls to the buffer manager.
 - **Idea:** Keep as many blocks (pages) in memory as possible to reduce disk accesses.
- **Replacement Policies**
 - e.g., **LRU** (Least-Recently-Used pages ... are discarded => the oldest are discarded first), **MRU** (Most-Recently-Used, the newest are discarded first), **LFU** (Least-Frequently-Used)
- **Prefetching or Double Buffering**
 - **Idea:** speed-up access by pre-loading “future”-needed data
 - **Cons:** requires extra main memory; no help if requests are random

Context of next slides



Alternative File Organizations

(Εναλλακτική Οργάνωση Αρχείων)



- **File organization (Οργάνωση Αρχείου):** Method for **arranging** a collection of records and supporting the concept of a **file**.
- In all file organizations the records are accessed by their respective **RecordID**
 - Note that a **Record ID (RID)** usually has the following structure (**PageID, SlotID**), where **SlotID** defines the offset : i) inside **PageID** at which RID begins; or ii) inside the Slot Directory that resides within page PageID (explained in next lectures) *Page i*
- **Basic Questions**
 - How to store data inside **data records** (fixed-length records vs. variable-length records)?
 - How to store data-records **inside a file** (heap file, sorted file, indexed file)?
 - How to make a certain File Organization more powerful by complementing them with an Index?



File Organization Types

(Τύποι Οργάνωσης Αρχείων)



- Heap files (Αρχεία Σωρού): Suitable when typical access is a **file scan (σάρωση αρχείου)** retrieving **ALL** records
 - Suitable for queries like “SELECT * FROM Employees;”
- Sorted Files (Ταξινομημένα Αρχεία): Best if records must be retrieved in some order, or only a **‘range’ (διάστημα)** of records is needed.
 - Suitable for queries like “SELECT * FROM Employees WHERE 20<age and age<30;”
- **Each file organization makes certain operations efficient, but we are interested in supporting more than one operation!**
- To deal with such situations the DBMS builds one or more **indexes**.
- An **index** on a file is designed to speed up operations that are **not efficiently** supported by the basic organization of that file.

Indexes (Access Methods)

(Ευρετήρια Δευτερεύουσας Μνήμης)



- An *index* is a data structure that has **index records** which **point to** certain **data records**.
- An index can **optimize** certain kinds of retrieval operations (depending on the index).

• Definitions



- **Index Page (Σελίδες Ευρετηρίου) vs. Data Pages (Σελίδες Δεδομένων):** Index Pages store index records to data records. Both reside on disk because we might have many of these pages!
- **Data Record (Εγγραφή Δεδομένων):** Stores the actual data e.g., (59,Mike,3.14) .
- **Index Record (Εγγραφή Ευρετηρίου):** Stores the RID of another index record or a data record.


B+ Tree Index Overview

(Σύνοψη του Ευρετηρίου B+ Tree)

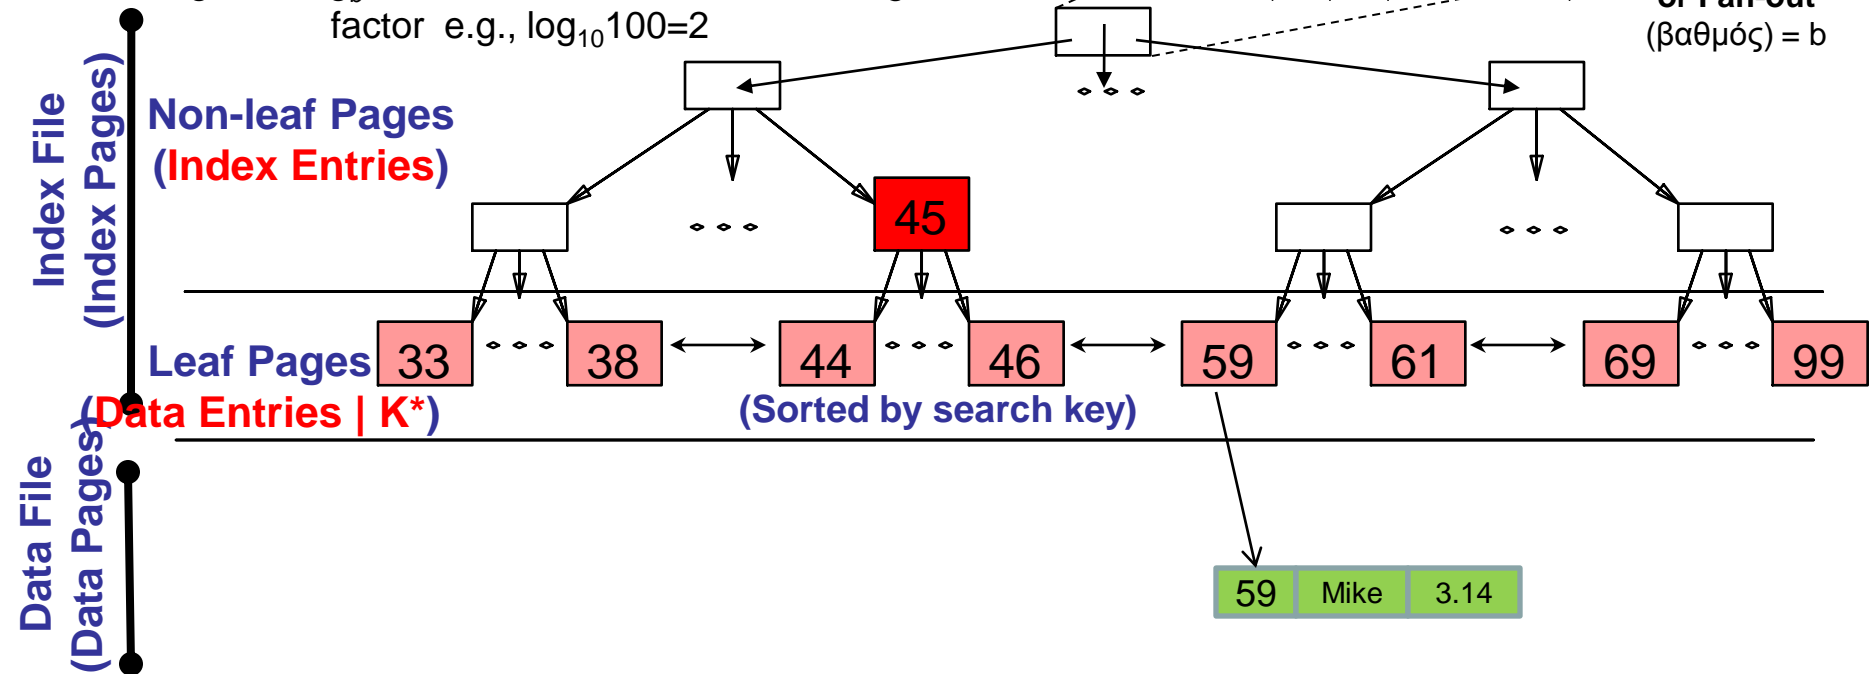


Balanced Tree (Ισοζυγισμένο Δένδρο)

Height = $\lceil \log_b N \rceil$, N :#data-entries, b =branching factor e.g., $\log_{10} 100 = 2$



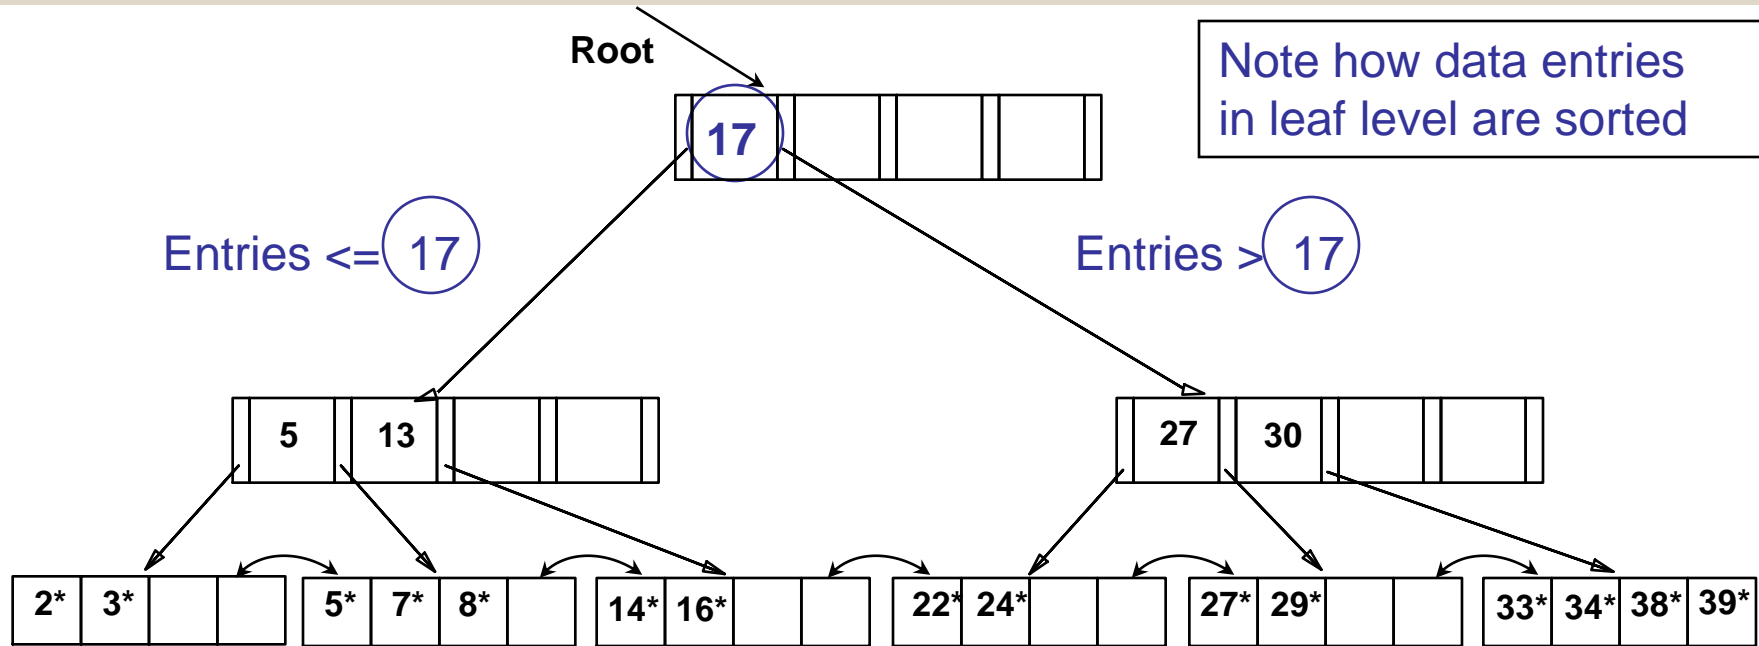
Branching Factor
or Fan-out
(βαθμός) = b



- ❖ Non-leaf pages have *index entries*; only used to direct searches.
- ❖ Leaf pages contain *data entries* K^* , and are chained (prev&next)
- ❖ The data records e.g., (59,Mike,3.14) could have been stored inside the respective data entry. Then the index file would be the same with the data file. ➔ **Index File Organization.**

Example B+ Tree

(Παράδειγμα Χρήσης B+ Tree)



- Find 28*? 29*? All $> 15^*$ and $< 30^*$
- Insert/delete: Find data entry in leaf, then change it. Need to adjust parent sometimes.
 - And change sometimes bubbles up the tree

Structure of Data Entry k^*

(Δομή της Καταχώρησης K^*)



- In a **data entry k^*** we can store:
 - **Alternative 1:** $\langle k \rangle$ (Key Value), or
 - **Alternative 2:** $\langle k, RID \rangle$ (Key Value, Data record with key value k), or
 - **Alternative 3:** $\langle k, [RID_1, RID_2, \dots, RID_n] \rangle$, where RID_i is a data record with key value k .
- *Choice of alternative for data entries is **orthogonal** to the indexing technique used to locate data entries with a given key value k .*
 - In particular, **ANY** of the above alternatives might be used with **ANY** index (hash or tree)

Data Entry k^* Examples

(Παραδείγματα Καταχώρησης k^*)



- **Alternative 1: $\langle k \rangle$**

Results in a
Index File Organization!

59, Mike, 3.14

Index Data Entry

- **Alternative 2: $\langle k, \text{RID} \rangle$**

59, RID#10

Index Data Entry

59 Mike 3.14

Data Record

RID#10

- **Alternative 3: $\langle k, [\text{RID}, \dots, \text{RID}] \rangle$**

59, RID#10, RID#61, #RID82

Index Data Entry

59 Mike 3.14

RID#10

59 Chris 33.14

RID#61

59 Jim 53.14

RID#82


Data Record

Clustered vs. Unclustered Indexes



(Ομαδοποιημένα vs. Μη-Ομαδοποιημένα Ευρετήρια)

- **Clustered Index (Ομαδοποιημένο Ευρετήριο):** If **order (διάταξη)** of **data records** is the same as, or 'close to', order of **data entries**, else called unclustered index.
- **Alternative 1** implies **clustered** (since datarec same as dateentry)
 - Faster Range Queries: Consecutive data records reside on the same page.
- **Alternatives 2,3** are usually **unclustered**.
 - Slower Range Queries: Consecutive data records reside on different pages



Lecture Outline


Overview of Storage and Indexing



- **Note:** The subsequent slides aim to ***qualitatively compare*** (**ποιοτική σύγκριση**) the file organization and indexes alternatives we introduced previously.

- 8.4) Comparison of File Organization

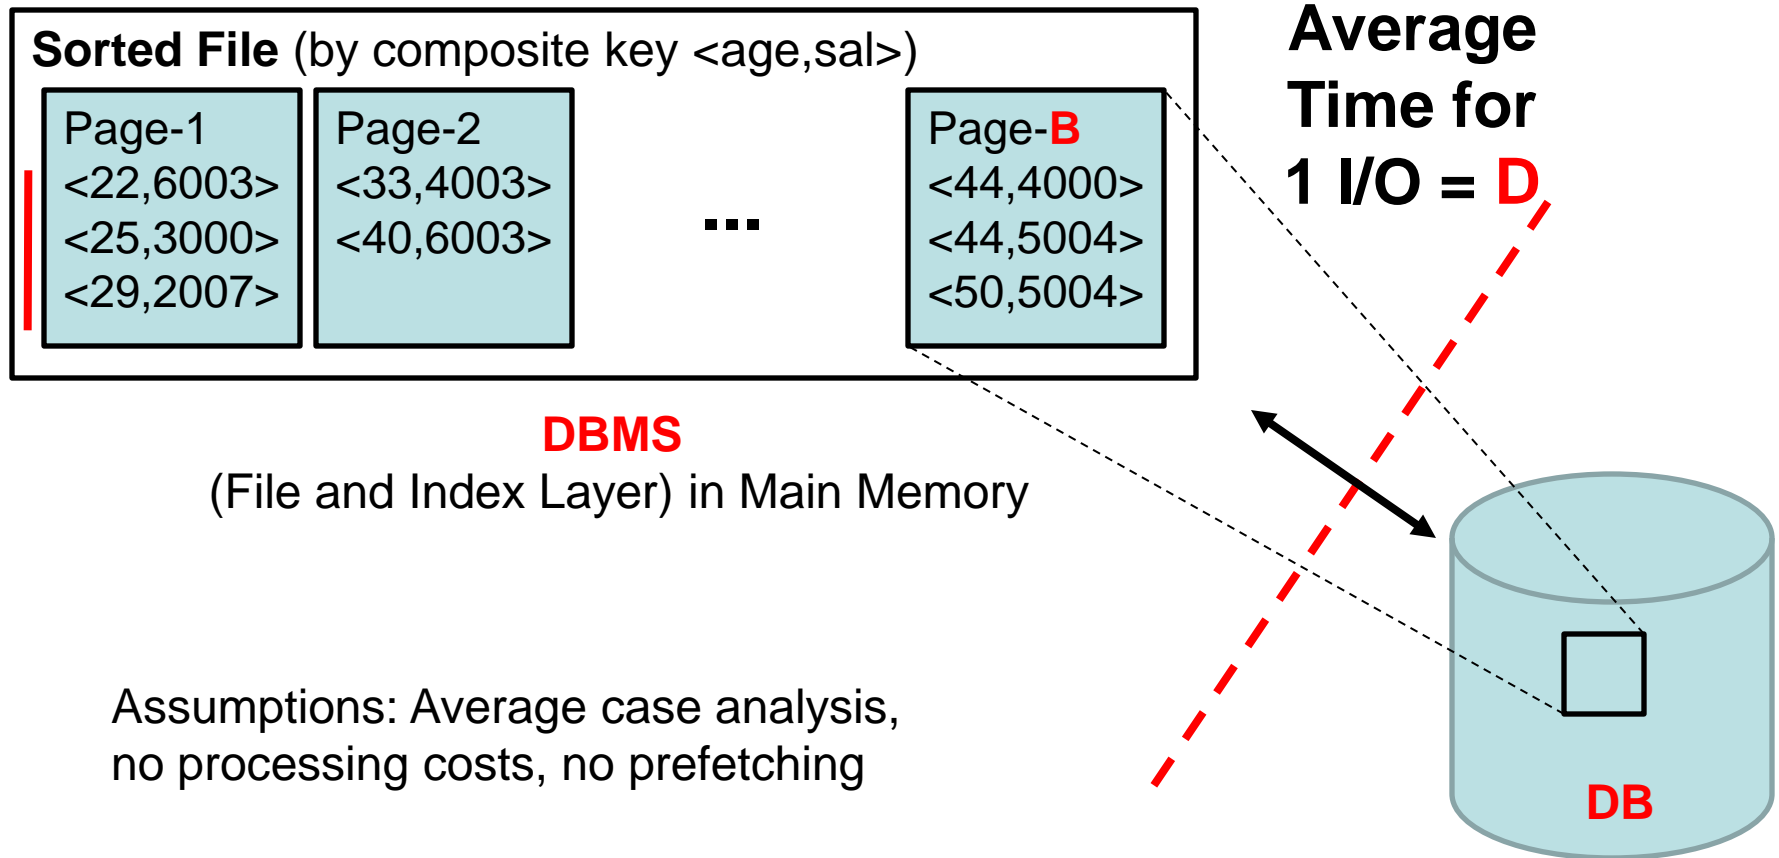
- System and Cost Model (Μοντέλο Κόστους)
- **Heap** Files, **Sorted** Files and **Clustered** Files
(Αρχεία: **Σωρού**, **Ταξινομημένα**, **Ομαδοποιημένα**)
- Comparison on I/O Costs (Σύγκριση Κόστους I/O)



- 8.5) Indexes and **Performance Tuning** (Ρύθμιση Επίδοσης)
 - Understanding the **Workload** (Εκτιμώντας τον Φόρτο Εργασίας)
 - Index **Specification** in SQL (Δήλωση Ευρετηρίων στην SQL)
 - **Index-Only Plans** (Πλάνα με Μόνο το Ευρετήριο)
 - Index Selection **Guidelines** (Οδηγίες Επιλογής Ευρετηρίων)

System Model

(Μοντέλο Συστήματος)



Operations to Compare

(Πράξεις που θα Συγκριθούν)



- **Scan (Σάρωση):** Fetch all records from disk
 - e.g., `SELECT * FROM Employees;`
- **Equality Selection (Επιλογή Ισότητας)**
 - e.g., `SELECT * FROM Employees WHERE age=33 AND sal=4003;`
- **Range selection (Επιλογή Διαστήματος)**
 - e.g., `SELECT * FROM Employees WHERE age BETWEEN 35 AND 45;`
 - e.g., `SELECT * FROM Employees WHERE 35<age AND sal<=4000;`
 - ***But NOT:** `SELECT * FROM Employees WHERE sal>40;` (tree index is on age ☹)*
- **Insert a record (Εισαγωγή Εγγραφής)**
 - e.g., `INSERT INTO Employees (age, sal) VALUES (45, 3000);`
- **Delete a record (Διαγραφή Εγγραφής)**
 - e.g., `DELETE FROM Employees WHERE age=45;`

Heap File Analysis

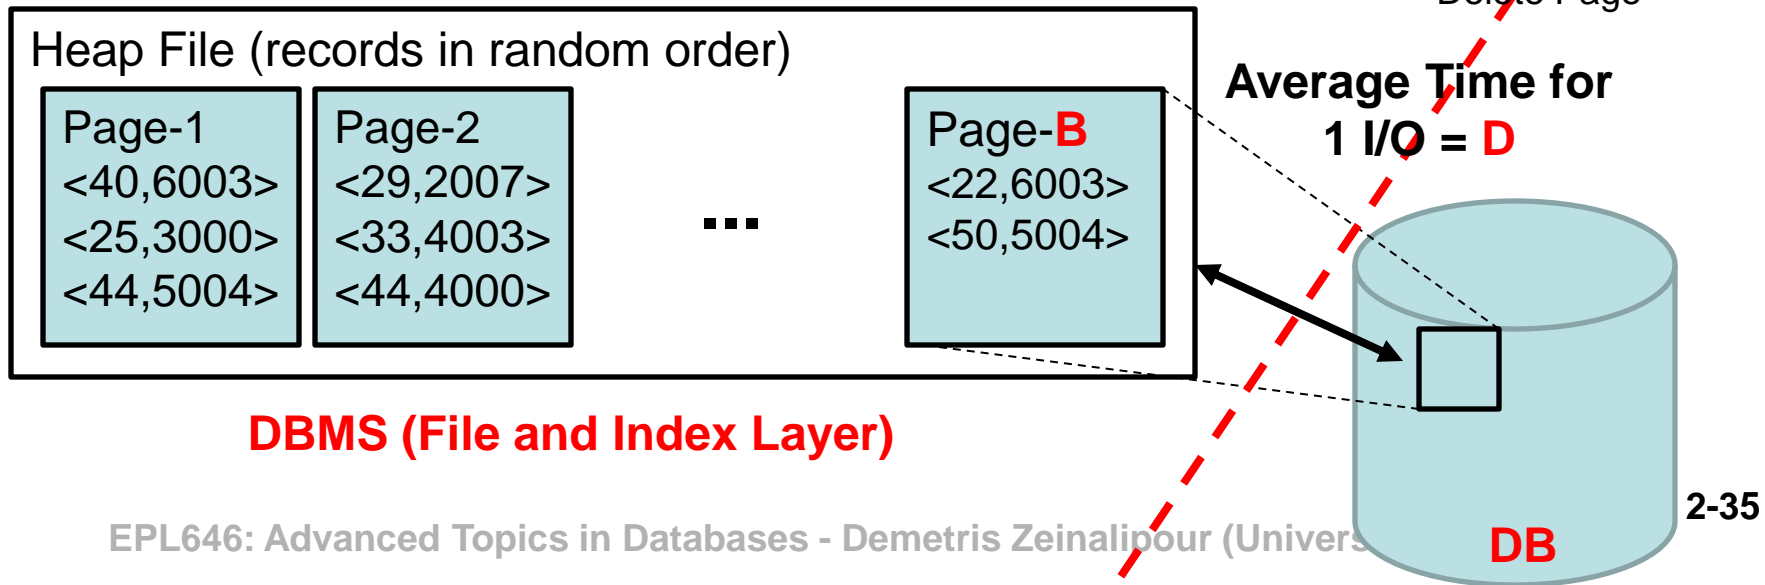
(Ανάλυση Αρχείου Σωρού)



- **Heap File Assumptions**


- Equality Selection on key <age,sal>
- Equality Selection produces exactly 1 match.

ScanAll	Eq. Selection	Range Selection	Insert	Delete
BD	0.5BD	BD	2D	0.5BD+ D
All records	On average we traverse ½ records	Traverse all to find pages in range	Read (last) PageB + Write PageB	Find Page + Delete Page



- Files compacted after deletions (no holes in pages)

All records	Binary Search over B pages. Each I/O costs D	Binary Search for 1 tuple, then transfer rest qualifying pages	Binary Search for Correct Position + Shift (Read/Write) $\frac{1}{2}$ subsequent pages (i.e., $2 \times 0.5BD = BD$)	Same as Insert but $\frac{1}{2}$ pages are shifted back in order to compact the file
-------------	--	--	---	--



Understanding the Workload

(Εκτιμώντας τον Φόρτο Εργασίας)



- **Workload (Φόρτος Εργασίας):** The typical mix of i) **Query (Select)** and ii) **Update (Insert/Delete/Update)** operations in a DBMS system.
- i) For each **query/update** in the **workload** :
 - Which types are involved (Select, Insert, Delete, Update)
 - Which **relations/attributes** (σχέσεις, χαρακτηριστικά) does it **access**?
 - Which **attributes** are involved in **selection/join** (επιλογή/ συνένωση) conditions? How **selective** are these conditions likely to be?

- **Selectivity (Επιλεκτικότητα της Συνθήκης):** The fraction of tuples selected by a selection condition is referred to as the selectivity of the condition.

E.g., $\sigma_{\text{age} > 40}(\text{EMPLOYEE})$ returns 10 out of 1000 tuples. **Selectivity=1%**

Index Specification in SQL



Δήλωση Ευρετηρίου στη SQL

- The **SQL standard (up until SQL 2008)** does not include any statement for **creating/dropping** indexes.
- **However**, in practice **every major DBMS supports such indexes (access methods)** such as **Btrees, Hash, Rtrees, GIST**.

Example from the PostgreSQL DBMS

```
CREATE INDEX AgeSalIndex  
  ON Employees (age, sal)  
  USING BTREE  
  WHERE sal > 3000
```

Choice of Indexes

(Επιλογή των Ευρετηρίων)



- The DBA is usually confronted with several questions in regards to indexes:
 - **Which relations** should have indexes?
 - What **type** of index should we use?
Clustered? Hash? Btree?
 - What **attribute(s)** should be the search key?
 - Should we build **several indexes**?

Index Selection Guidelines

(Οδηγίες Επιλογής των Ευρετηρίων)



- **Tip 1:** Consider the **queries executed most of the time** (most important ones), e.g., for Oracle :
 - SELECT executions, sql_text FROM v\$sqlarea ORDER BY executions desc;
 - V\$ => Oracle's Dynamic Performance Views
- **Tip 2:** Try to choose indexes that **benefit as many queries as possible**
- **Tip 3:** Attributes in **WHERE** clause are candidates for index keys.
- **Tip 4: Hash vs. Tree**
 - **Exact match** condition suggests **Hash index**.
 - **Range query** suggests **tree index**.

Index Selection Guidelines

(Οδηγίες Επιλογής των Ευρετηρίων)



- **Tip 5:** Consider the best plan using the **current indexes**, and see if a better plan is possible with an additional index. If so, create it!
- **Tip 6:** Since only one index can be **clustered** per relation, choose it based on important queries that would benefit the most from clustering.
- **Tip 7: Multi-attribute $\langle a, b, c \rangle$ search keys** should be considered when a WHERE clause contains several conditions (e.g., $a=3$ and $b>3$ and $c>3$).
- **Tip 8:** Indexes can make **queries go faster** but **updates become slower**. Indexes also require additional disk space, choose them wisely!