



EPL646 – Advanced Topics in Databases

Lecture 13

NoSQL Databases: CouchDB II (Semi-structured JSON DB)

**Chapter 20: Abiteboul et. Al.
+ <http://guide.couchdb.org/>**

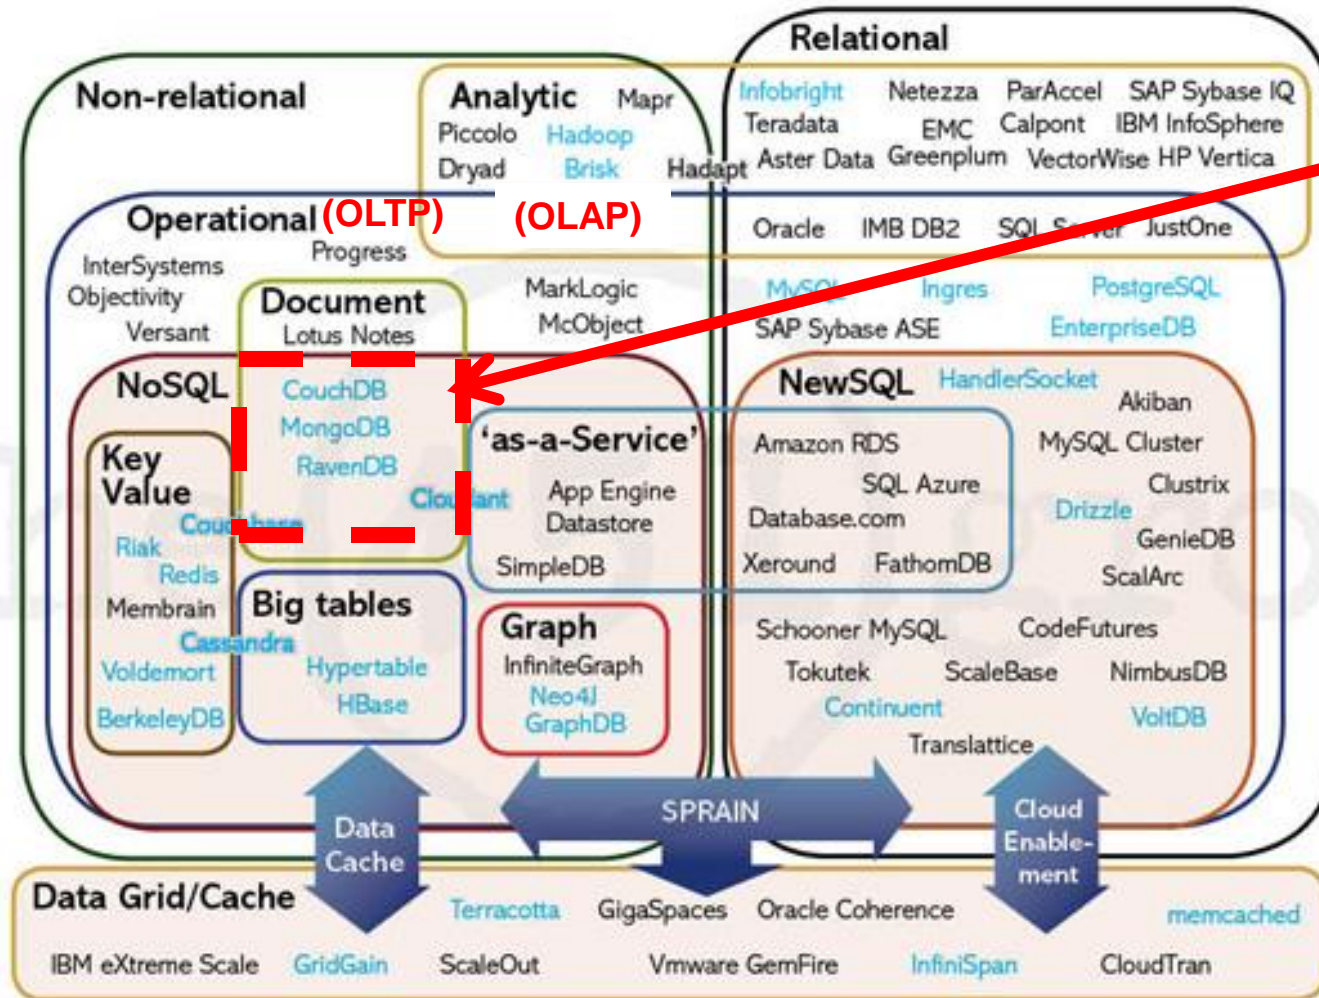
Demetris Zeinalipour

<http://www.cs.ucy.ac.cy/~dzeina/courses/epl646>

EPL646: Part B



Distributed/Web/Cloud DBs/Dstores



Lecture
Focus

Venn
Diagram by
451 group

<http://xeround.com/blog/2011/04/newsq-cloud-database-as-a-service>

Lecture Outline

(Introduction to Semi-structured Data)

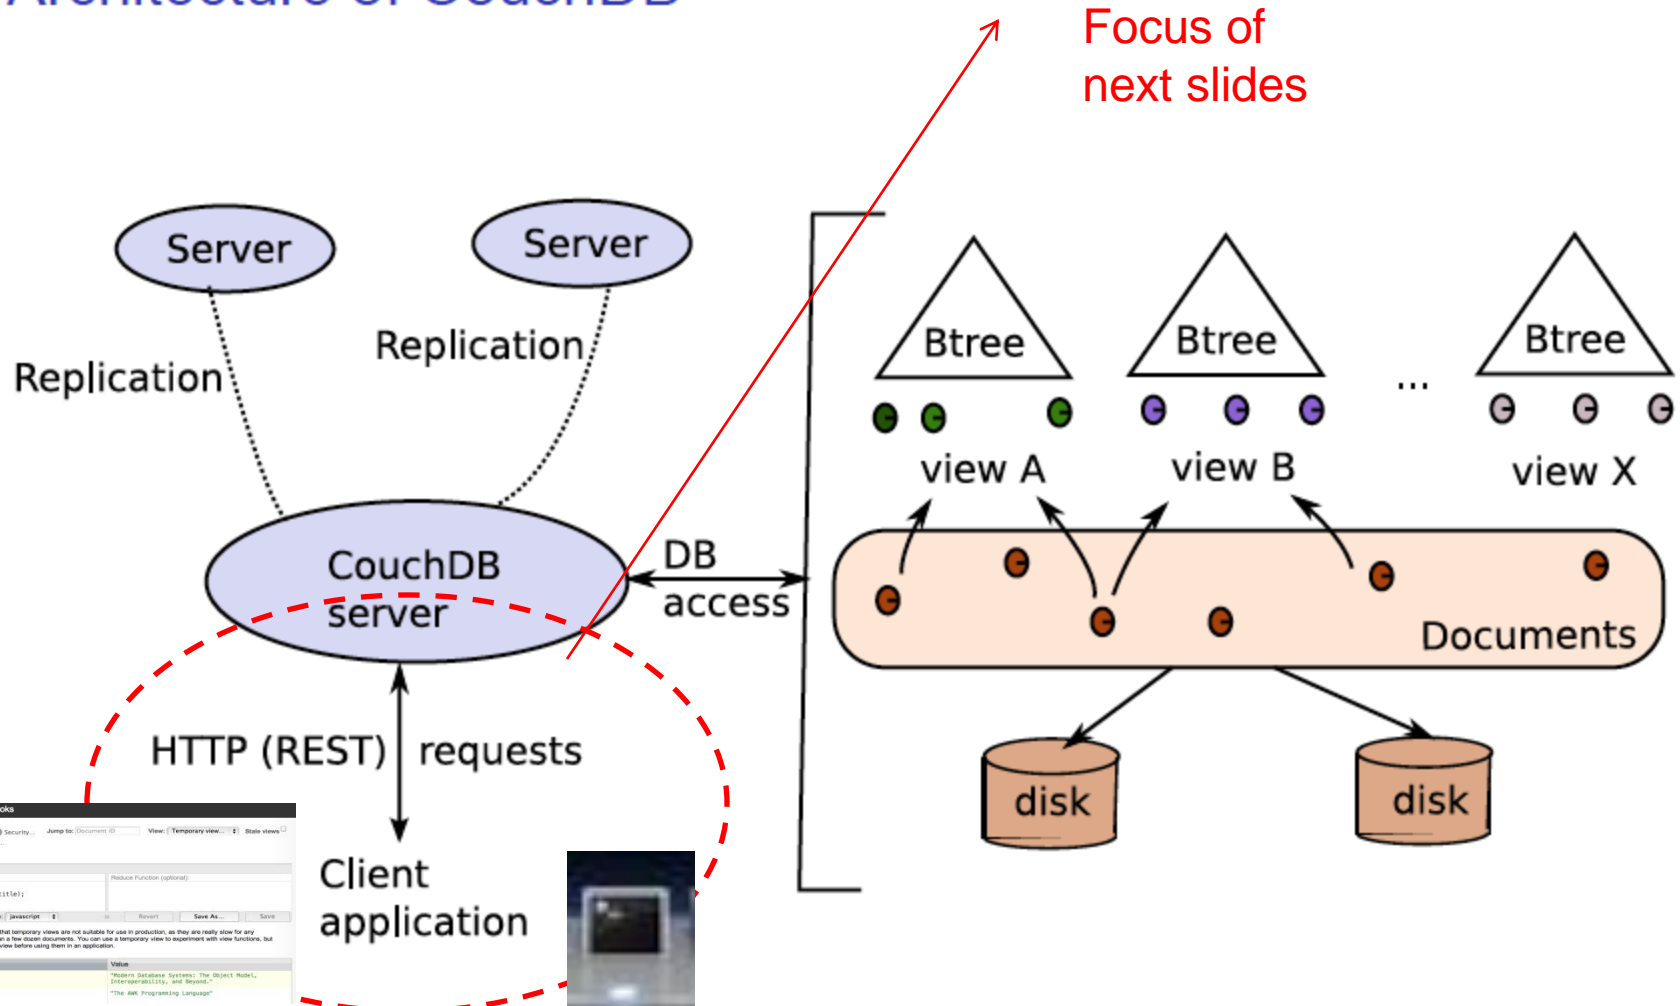


- Intro to Web2.0 & JSON Data Interchange Format
- JSON Key-Value Data Model
- CouchDB: A JSON Database (written in Erlang)
 - Using Command Line CURL/ Web-based FUTON
 - CouchDB Architecture (Btrees, Filesystem, Replication)
 - REST Principles
 - Creating DBs, Adding Docs, Updating Docs, Deleting Docs, _ID and _REV issues, Multi-Version CC (MVCC)
 - Querying Data with (Materialized) Views (Map-Reduce style in Javascript)
 - Replication and Scalability Issues

CouchDB REST Interface



Architecture of CouchDB



REST Principles



Apart©: REST principles

Roy Fielding (Univ. of Calif, Irvine PhD)

Founder of Apache HTTP Project
HTTP 1.0 (RFC1945) w/ Berners-Lee
HTTP 1.1 (RFC2616) w/ others



A Web-service dialect that enables exchanges of HTTP messages to access, create, and manage resources.
protocol as follows:

GET retrieves the resource referenced by the URI.

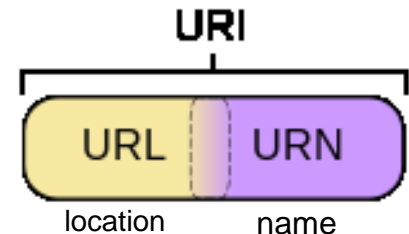
PUT creates the resource at the given URI.

POST sends a message (along with some data) to an existing resource.

DELETE deletes the resource.

GET

```
/test/demo_form.asp?name1=value1&name2=value2
```



POST

```
POST /test/demo_form.asp HTTP/1.1  
Host: w3schools.com  
name1=value1&name2=value2
```

Very convenient in a Web environment: no need to use a client library –
Documents can easily be incorporated in a Web interface.

CouchDB: CREATE DB / INSERT (with CURL)



A short interactive session

Curl: Client URL Command Line Tool

Talk to the server: send an HTTP request, get a response.

```
$ curl -X GET http://mycouch.org  
{ "couchdb": "Welcome", "version": "1.0.1" }
```

Create a db = put a resource (the name suffices).

```
$ curl -X PUT http://mycouch.org/myDB  
{ "ok": true }
```

Create a document = put a resource in a db (give the JSON document in the HTTP request).

```
$ curl -X PUT http://mycouch.org/myDB/myDoc \  
      -d '{ "key": "value" }'  
{ "ok": true, "id": "myDoc", "rev": "1-25eca" }
```

Get the document after its URI:

```
$ curl -X GET http://mycouch.org/myDB/myDoc  
{ "_id": "myDoc", "_rev": "1-25eca", "key": "value" }
```

Retrieve all tables (e.g., show tables; in MySQL)

```
$ curl -X GET http://127.0.0.1:5984/_all_dbs  
[ "_replicator", "_users", "books", "booksreplica", "movies", "twitter" ]
```

CouchDB: _ID and _REV



Document management in CouchDB

Field	Value
_id	"book1.json"
_rev	"1-410c67caca526b476abc72e73b003605"

Each document has an **id** and a **revision number**.

Use: `uuidgen` to generate uniq ids during insert

UUID=128bit (32Hex digits) =
 2.4×10^{38} keys

Each update to a document creates a new version, with the same **_id** but a new revision number.

Validation functions can be assigned to a collection: any document inserted or updated must be validated by these functions (ad-hoc type-checking). *Like triggers...*

Design documents are a special type of CouchDB document that contains application code.

A **view** is a new key-document collection, specified via MAPREDUCE.

(we will see views later)

Documents can be **replicated** in other CouchDB instances.

(we will see replication later)

CouchDB: BULK LOAD (with CURL)



```
# Download the files from the web
#(or wget http://webdam.inria.fr/Jorge/files/jsonmovies.zip)
$ curl -O http://webdam.inria.fr/Jorge/files/jsonmovies.zip

# Unzip Movies
$ unzip jsonmovies.zip

# or assign unique IDs
$COUCHDB/movies/`uuidgen`

# List the files
$ ls -al | head
total 12480
-rw-r--r--@    1 dzeina  staff    218   9 ??t   2011 book1.json
-rw-r--r--@    1 dzeina  staff    222   9 ??t   2011 book10.json
-rw-r--r--@    1 dzeina  staff    197   9 ??t   2011 book100.json

# Bulk load using Bash
$ for i in `ls .`; do curl -X PUT $COUCHDB/movies/$i -d \@$i; done
{"ok":true,"id":"book1.json","rev":"1-410c67caca526b476abc72e73b003605"}
{"ok":true,"id":"book10.json","rev":"1-d0cc2ae0ab3211314a65a5c5244df221"}
{"ok":true,"id":"book100.json","rev":"1-2cfe83eea8cad920cfd66755ac78b46f"}
```


CouchDB: UPDATE / DELETE (with CURL)



Updating data

Updating in COUCHDB = adding a new version.

COUCHDB applies a **Multi-version concurrency control** protocol which requires that you send the version that must be updated:

```
$ curl -X PUT $IP/movies/tsn?rev=1-db1261 -d @newDoc.json \  
-H "Content-Type: image/jpg" \  
{ "ok":true, "id": "tsn", "rev": "2-26863" }
```

```
{  
  "_attachments": {  
    "d1.tiff": {  
      "content_type": "image/tiff",  
      "revpos": 6,  
      "digest": "md5-2SYLdVnPnGNXPz10As8X5g==",  
      "length": 123776,  
      "stub": true  
    }  
  }  
}
```

Image separated from json

Deletion is obtained with DELETE.

```
$ curl -X DELETE $COUCHADDRESS/movies/tsn?rev=2-26863 \  
{ "ok":true, "id": "tsn", "rev": "3-48e92b" }
```

A new version has been created! (logical deletion).

CouchDB: UPDATE / DELETE (MVCC Explained)



- **Multi-Version CC (MVCC)** uses timestamps or increasing IDs to achieve transactional consistency.
 - MVCC provides each user with a snapshot of the database
- The database holds multiple versions (e.g., t0, t1 for Obj1)

Time	Object 1	Object 2	
t1	"Hello"	"Bar"	Latest Version
t0	"Foo"	"Bar"	

- Assume that a **T1** is a long-running READ xact started with state "t1"
- Also assume that **T2** is a WRITE transaction that creates state "t2"

T1 continues to have access to the t1 state! (thus, the necessary ISOLATION)

Time	Object 1	Object 2	Object 3	
t2	"Hello"	(deleted)	"Foo-Bar"	Latest Version
t1	"Hello"	"Bar"		
t0	"Foo"	"Bar"		

For **serializability / recoverability**: Multiversion Histories need to follow similar ideas to Single-version Histories we saw with Basic TO and Strict TO.

CouchDB: UPDATE / DELETE (with CURL)



```
# UPDATE document "4C4F2E4C-E1AC-4C80-B90E-A9D0BCB638C8"
$ curl -X PUT $COUCHDB/googlebooks/4C4F2E4C-E1AC-4C80-B90E-
A9D0BCB638C8?rev=1-1d974c2aadd42b0b8699678d61a0e4ff -d
@potter.json

{"ok":true,"id":"4C4F2E4C-E1AC-4C80-B90E-
A9D0BCB638C8","rev":"2-40d4b6bf3530e6af3a84904652ce9a8c"}

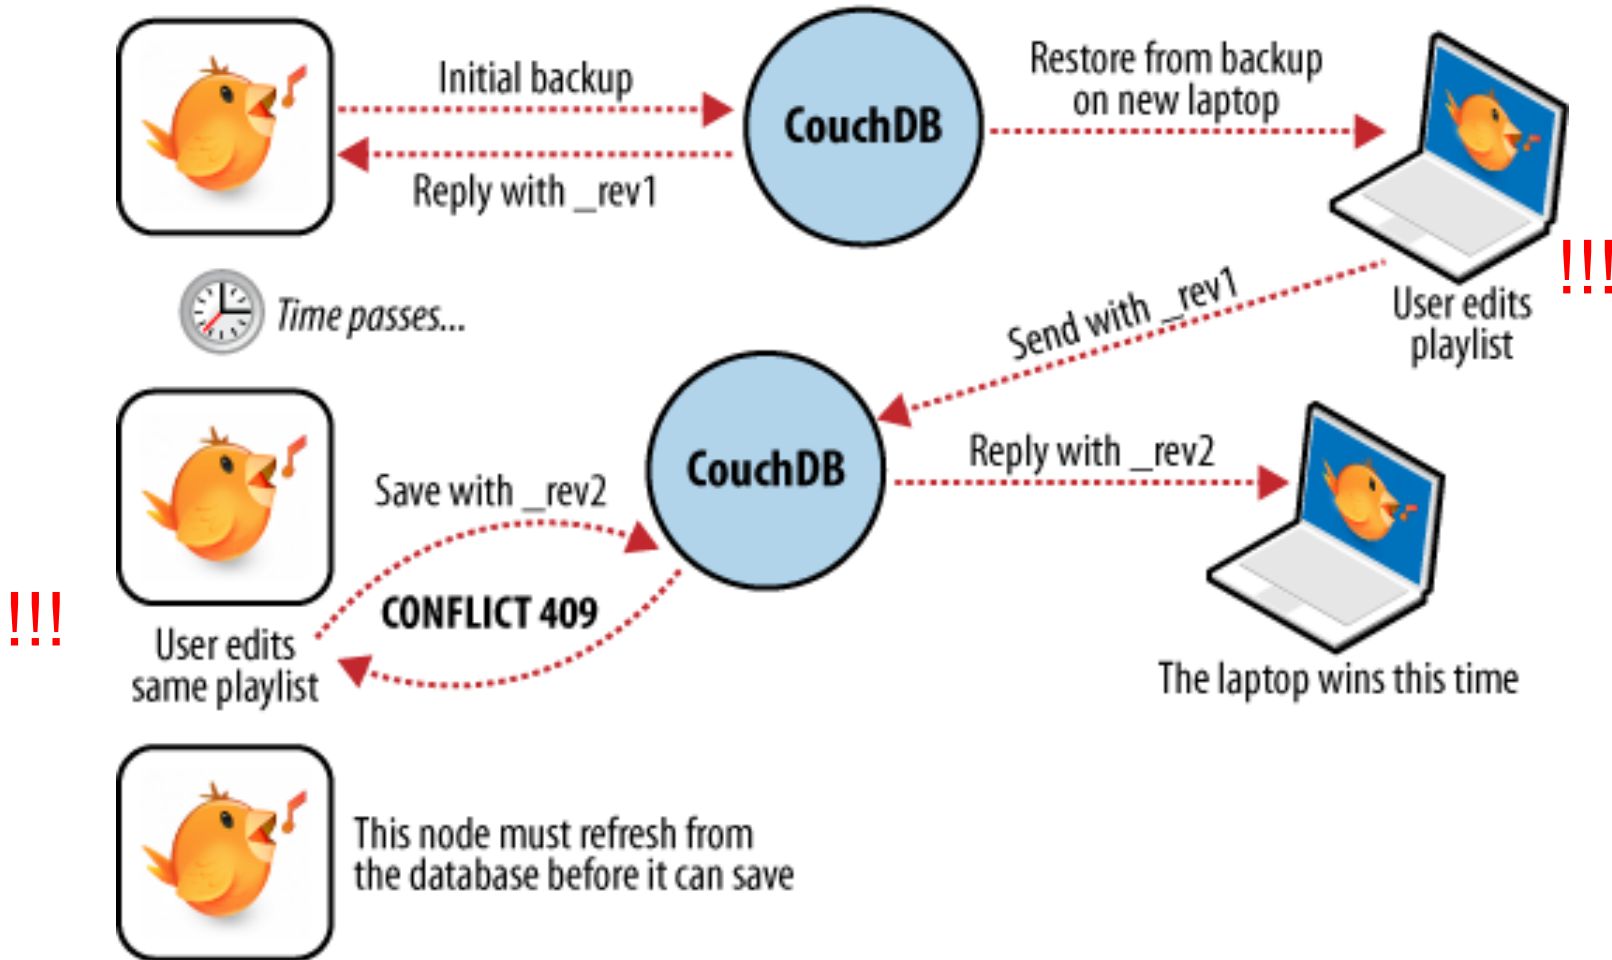
# Now DELETE PRELAST version "rev=1-
1d974c2aadd42b0b8699678d61a0e4ff"
$ curl -X DELETE $COUCHDB/googlebooks/4C4F2E4C-E1AC-4C80-B90E-
A9D0BCB638C8?rev=1-1d974c2aadd42b0b8699678d61a0e4ff
{"error":"conflict","reason":"Document update conflict."}

# DELETE LAST version "rev=2-40d4b6bf3530e6af3a84904652ce9a8c"
$ curl -X DELETE $COUCHDB/googlebooks/4C4F2E4C-E1AC-4C80-B90E-
A9D0BCB638C8?rev=2-40d4b6bf3530e6af3a84904652ce9a8c
{"ok":true,"id":"4C4F2E4C-E1AC-4C80-B90E-
A9D0BCB638C8","rev":"3-a77d6120602b51cbf9b8663c1ee8f9e3"}
```

CouchDB: UPDATE / DELETE (Example)



Syncing playlists between multiple Songbird clients



CouchDB: SELECT QUERY (expressed as Javascript)



Views in COUCHDB

A **view** is the result of a MAPREDUCE job = a list of (*key, value*) pairs.

- + Temporary View: executed on demand (e.g., through Futon) – good for development
- + Permanent (Materialized) Views: also called design documents, accessible through URI

Views are **materialized** and indexed on the key by a B+tree.

(we will study the Map-Reduce Programming Model more extensively in the next lecture)

A MAP function

```
function(doc)
{
    emit(doc.title, doc.director)
}
```

Similar to SQL
SELECT fields

A REDUCE function

```
function (key, values) {
    return values.length;
}
```

Similar to SQL GROUP BY
AGGREGATES;

Lecture Outline

(Introduction to Semi-structured Data)



- **SQL: SELECT * FROM Books;**

```
function(doc) {  
    emit(null, doc);  
}
```

Key ▲	Value
null ID: book1.json	{_id: "book1.json", _rev: "1-410c67caca526b476abc72e73b003605", type: "Book", title: "Modern Database Systems: The Object Model, Interoperability, and Beyond.", year: "1995", publisher: "ACM Press and Addison-Wesley", authors: [], source: "DBLP"}
null ID: book10.json	{_id: "book10.json", _rev: "1-d0cc2ae0ab3211314a65a5c5244df221", type: "Book", title: "The AWK Programming Language", year: "1988", publisher: "Addison-Wesley", authors: ["Alfred V. Aho", "Brian W. Kernighan", "Peter J.

- **SQL: SELECT pub FROM Books;**

```
function(doc) {  
    emit(doc._id, doc.publisher);  
}
```

Key ▲	Value
"book1.json" ID: book1.json	"ACM Press and Addison-Wesley"
"book10.json" ID: book10.json	"Addison-Wesley"

*ID always part
of answer (but
not _REV)*

- **SQL: SELECT pub, typ FROM Books**

```
function(doc) {  
    emit(doc._id,  
        {"pub": doc.publisher, "typ": doc.type});  
}
```

Key ▲	Value
"book1.json" ID: book1.json	{pub: "ACM Press and Addison-Wesley", typ: "Book"}
"book10.json" ID: book10.json	{pub: "Addison-Wesley", typ: "Book"}

CouchDB: SELECT QUERY (with FUTON)



SQL Equivalent: **SELECT B.title FROM Books B;**

Overview > books

+ New Document ? Security... Jump to: View: **Temporary view...** Stale views ☐

⌵ Compact & Cleanup...
✕ Delete Database...

View Code **MAP (Query)**

Map Function:

```
function(doc) {  
  emit(doc._id, doc.title);  
}
```

Reduce Function (optional):

Run Language: javascript = Revert Save As... Save

Warning: Please note that temporary views are not suitable for use in production, as they are really slow for any database with more than a few dozen documents. You can use a temporary view to experiment with view functions, but switch to a permanent view before using them in an application.

Results

Key	Value
"book1.json" ID: book1.json	"Modern Database Systems: The Object Model, Interoperability, and Beyond."
"book10.json" ID: book10.json	"The AWK Programming Language"

Not saved yet => No
Design document =>
no Index yet

CouchDB: SELECT-WHERE QUERY (with FUTON)



SQL Equivalent: **SELECT B.title FROM Books B WHERE B.publisher="AW";**

⊕ New Document ⓘ Security... Jump to: View: **Temporary view...** Stale views ☐

⊗ Compact & Cleanup...

⊗ Delete Database...

▼ View Code

Map Function:	Reduce Function (optional):
<pre>function(doc) { if (doc.publisher == "Addison-Wesley") { emit(doc._id, doc.title); } }</pre>	

Run Language: **javascript** = Revert Save As... Save

Warning: Please note that temporary views are not suitable for use in production, as they are really slow for any database with more than a few dozen documents. You can use a temporary view to experiment with view functions, but switch to a permanent view before using them in an application.

Results

Key ▲	Value
"book10.json" ID: book10.json	"The AWK Programming Language"
"book100.json"	"The Design and Analysis of Spatial Data Structures"

CouchDB: Array Object Iteration (with FUTON)



```
{ Document
  "_id": "book10.json",
  "_rev": "1-d0cc2ae0ab3211314a65a5c5244df221",
  "type": "Book",
  "title": "The AWK Programming Language",
  "year": "1988",
  "publisher": "Addison-Wesley",
  "authors": [
    "Alfred V. Aho",
    "Brian W. Kernighan",
    "Peter J. Weinberger"
  ],
  "source": "DBLP"
}
```

Map Function

```
function(doc) {
  for (i in doc.authors) {
    author = doc.authors[i];
    emit(doc._id, author);
  }
}
```

Check "View Cookbook for SQL Jockeys" for more!
<http://guide.couchdb.org/editions/1/en/cookbook.html>

Results

Key ▲	Grouping: exact ▴	Value
"book10.json" ID: book10.json		"Alfred V. Aho"
"book10.json" ID: book10.json		"Brian W. Kernighan"
"book10.json" ID: book10.json		"Peter J. Weinberger"

CouchDB: Array Object Iteration (with FUTON + CURL)



Accessing views

Here is a view (without reduce function).

Array object iterator

```
function(doc)
{
  for (i in doc.actors) {
    actor = doc.actors[i];
    emit({"fn": actor.first_name, "ln": actor.last_name},
    }
  }
}
```

Save it in the **design document** named `examples`, and name the view `actors`. The view can be queried with:

```
$ curl $IP/movies/_design/examples/_view/actors
{"total_rows":16,"offset":0,
"rows":[
  {"id":"bed7",
   "key":{"fn":"Andrew","ln":"Garfield"},"value":"The Socia
  {"id":"91631b",
   "key":{"fn":"Clint","ln":"Eastwood"},"value":"Unforgiven
```

Result

EI

Like XSLT

CouchDB: Design Documents (Apps are Documents!)



CouchDB Design Documents can lead to "Standalone Web Apps"

Swinger New Play Export

Create a swinger' presentation

[Go!](#)

Your Presentations

BET HA BRACHA
Quality Products from Israel

Home About Us Contact Us

Tel: 01353 662981

Ely Service

Specialists in Classic Rovers

We specialise in classic Rovers (P4, P5 and P6 models) as well as other classic cars, ranging from a 1937 Austin 10 Cambridge through the decades to modern-day classics from the '90s.

Classic Car Services

- Basic servicing
- M.O.T.
- complete restoration, including:
 - mechanical work
 - bodywork
 - paintwork
 - interior trim
 - rechroming
 - supply and fitting of correct tyres for your classic.

[Read more...](#)

Modern Car Services

For those of you unlucky enough ever to have anything but a modern 'tin box'—but we love them really!—we offer you a similar level of service to that which our classic-car customers enjoy.

[Read more...](#)

History

The business first opened sometime in 1930, but the important date is March 1st, 1931, when Mr Charles Hull purchased the business. The following years saw the Company go from strength to strength, enjoying a

A Message from the Owner

I bought my first P6 in 1990 and they have gradually taken over. Over time I then progressed into P5 ownership as well. So when my chance came to take over at Ely, it seemed like manna from heaven!

July 2009

today < >

Sun	Mon	Tue	Wed	Thu	Fri	Sat
28 3a Amsterdam.rb	29	30 12a Kingsofcode	1 12a Kingsofcode 5a SysMap	2 6a Add new calendar item 9p Armageddon Day	3	4 6a foo 3p asf
5	6 11p hello	7	8 11a ruleza	9	10	11

[Help](#) | [Shopping Cart](#)

[Search](#)



Blue Kippah with Khaki border
£2.99



White Kippah with Black and Blue border
£2.99



White Kippah
£2.99



Messianic Scripture Songs from Israel
£8.55



The Jewish Album - Jonathan Settel
£8.55



Silver-Plated Jerusalem Talit Clip
£9.99

More Information

[Shipping & Returns](#)
[Privacy Policy](#)
[Conditions of Use](#)
[Links & Resources](#)
[How to Order](#)
[RSS Feeds](#)

Currency

Pound Sterling (£)

Best Sellers

Acrylic Prayer Shawl, Blue & Silver (23"x70")
Blue Kippah with Khaki border
White Kippah with Black and Blue border

1a Corp Conzo
9p Super 64 Tournament

CouchDB: Querying Materialized Views (with FUTON)



Querying views (Permanent | Materialized)

A view is a B+tree index. So:

```
function(doc)
{
  emit(doc.genre, doc.title) ;
}
```

is equivalent to

```
create index on movies (genre);
```

Recall the B+trees support key and range queries:

```
$ curl $IP/movies/_design/examples/_view/genre?key=\"Drama\"
{"total_rows":5,"offset":2,"rows":[
{"id":"9163", "key":"Drama","value":"Marie Antoinette"},
{"id":"bed7", "key":"Drama","value":"The Social network"}
]}
```

For range queries, send the two parameters `startkey` and `endkey`. (next slide)

CouchDB: Range Queries (with CURL)



CURL Command

```
curl -X GET  
http://127.0.0.1:5984/books/_design/authors/_view/autho  
rs?startkey=\"book980.json\"&endkey=\"books998.json\"
```

Results

View count

First result occurrence

```
$ {"total_rows":1893,"offset":1871,"rows":[  
{"id":"book980.json","key":"book980.json","value":"A. J. Kfoury"},  
{"id":"book980.json","key":"book980.json","value":"Michael A. Arbib"},  
{"id":"book980.json","key":"book980.json","value":"Robert N. Moll"},  
{"id":"book981.json","key":"book981.json","value":"Peter D. Mosses"},  
{"id":"book982.json","key":"book982.json","value":"Anne Mulkers"},  
{"id":"book983.json","key":"book983.json","value":"ller"},  
{"id":"book984.json","key":"book984.json","value":"ller"},
```


CouchDB Querying with Aggregation (Map & Reduce)



90% of queries will be MAP queries.

REDUCE allows us to carry out aggregation on a given field.

SQL:
SELECT key, sum(key)
FROM Books
GROUP BY key

Overview > books

+ New Document ? Security... Jump to: View: Stale views ☐

⊗ Compact & Cleanup...
⊗ Delete Database...

▼ View Code _design/examples

Map Function:	Reduce Function (optional):
<pre>function(doc) { emit(doc.publisher, 1); }</pre>	<pre>function(key, values, rereduce) { return sum(values); }</pre>

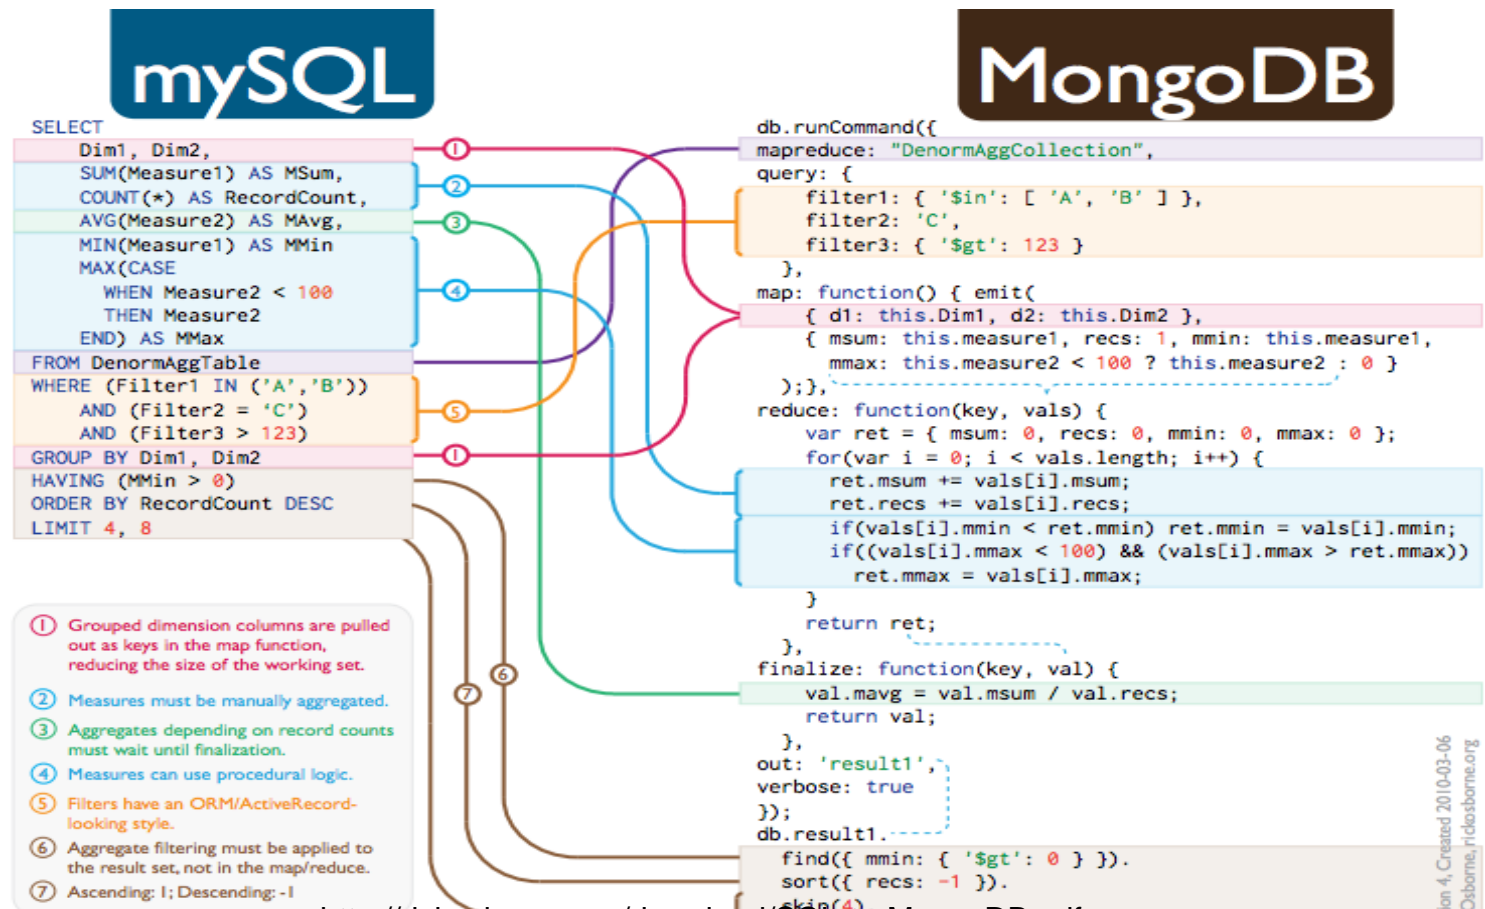
Run Language: = Revert Save As... Save

Key ▲	Grouping: <input type="text" value="exact"/>	Value	<input checked="" type="checkbox"/> Reduce
" ACM Press"		1	
" Claypool Publishers"		83	
" Francis"		1	
" Row"		1	
" Sons"		1	



Complex Map & Reduce

For a real app we could envision much more complex queries.



<http://rickosborne.org/download/SQL-to-MongoDB.pdf>

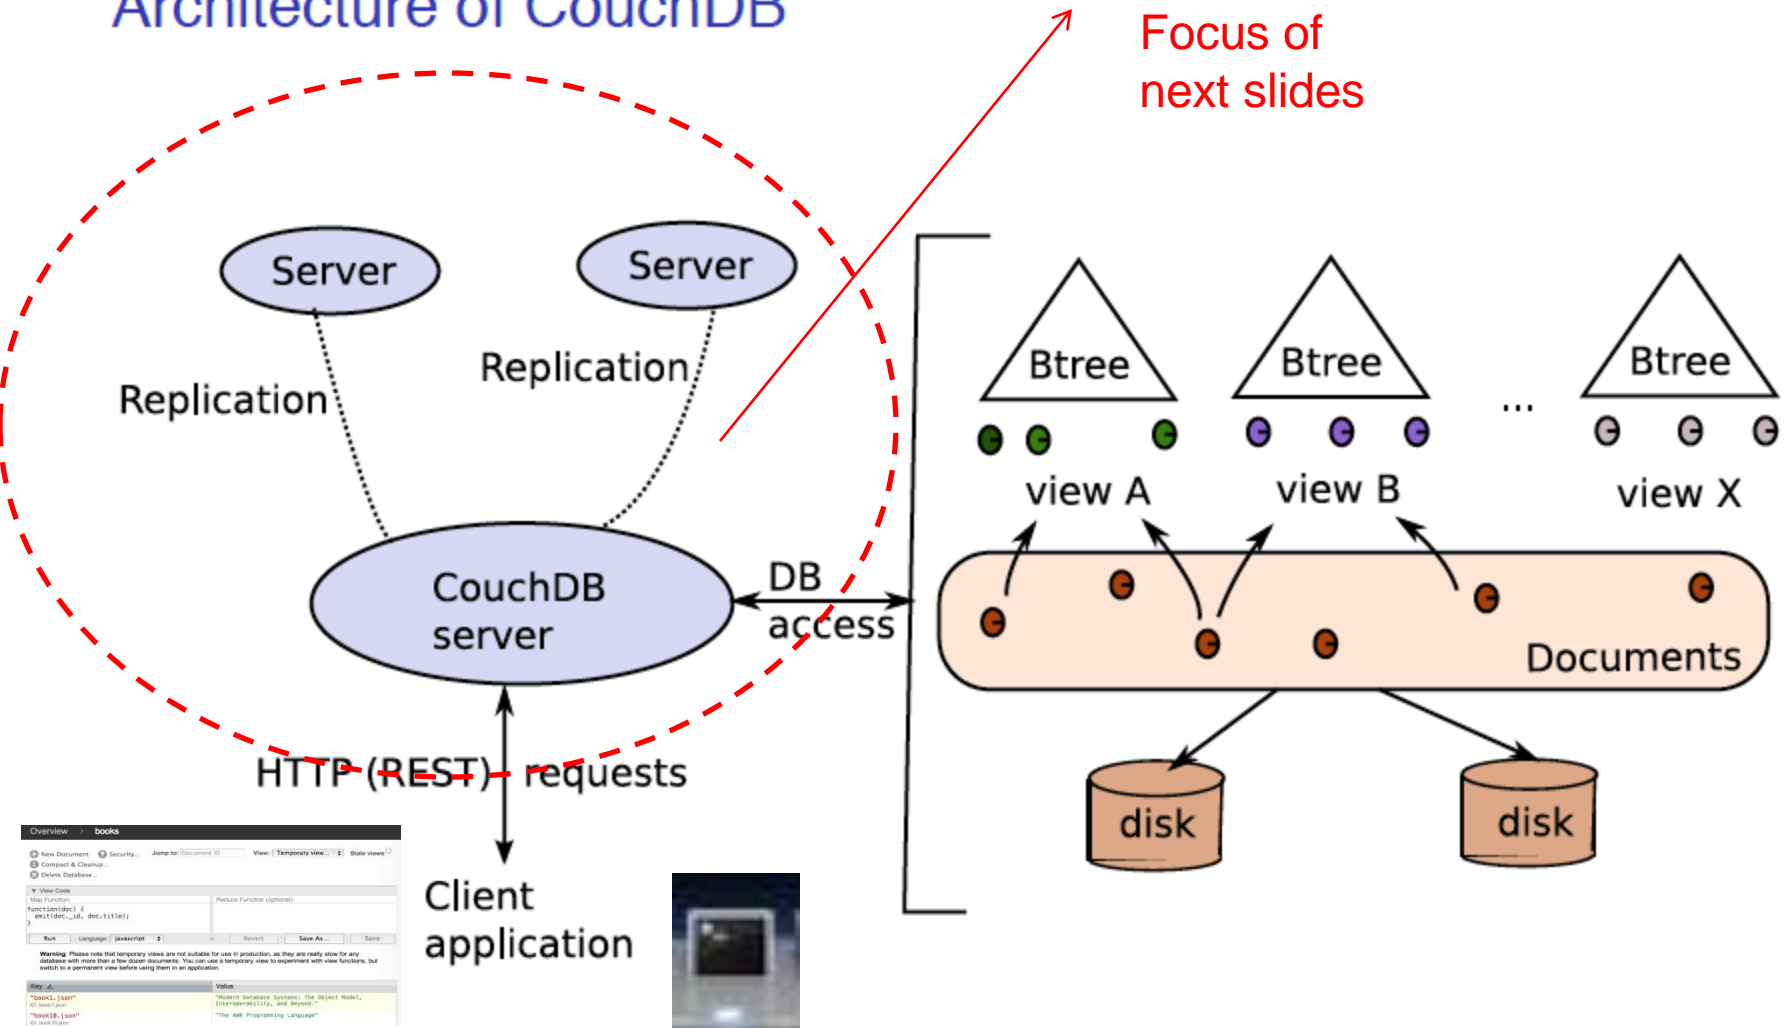
ion 4, Created 2010-03-06
Osborne, rickosborne.org



CouchDB Replication

Architecture of CouchDB

Focus of
next slides



CouchDB One-way/Symmetric Replication (with CURI)



The replication primitive

COUCHDB supports natively one-way **replication** from one instance to another.

```
curl -X POST $COUCHADDRESS/_replicate \  
  -d '{"source": "movies", "target": "backup",  
      "continuous": true}' \  
  -H "Content-Type: application/json"
```

That's all: any change in `movies` is automatically reported in `backup`.

Two way replication can be achieved by executing the inverse statement as well.

CouchDB One-way/Symmetric Replication (with CURL)



[Overview](#) > [Replicator](#)

Replicate changes from:

☒ Local Database:

☐ Remote database:

⇔

to:

☒ Local database:

☐ Remote database:

☒ Continuous

Event

No replication

CouchD
rel

Tools

- Overview
- Configuration
- Replicator**
- Status
- Verify Installation

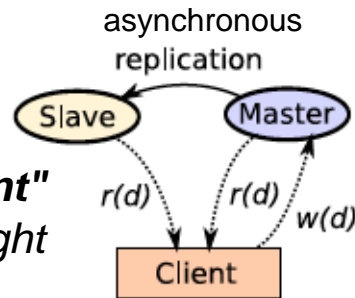
CouchDB Distribution (Concepts)



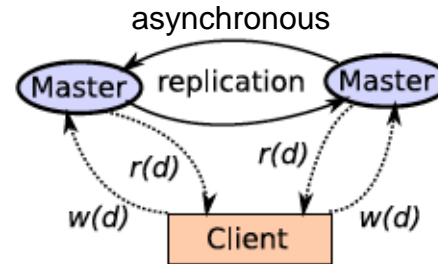
Distribution strategies

Combine a proxy that distributes requests, with the replication feature of COUCHDB.

"Eventually Consistent"
(reading from slave might be outdated)



a - Master-slave arch.



b - Master-master arch.

"Conflict Management is Necessary"

```
function(doc) {  
  if (doc._conflicts) {  
    emit(doc._conflicts, null);  
  }  
}
```

CouchDB Security



- CouchDB allows any request to be made by anyone (i.e., by default admin).
 - By default, CouchDB will listen only on your loopback network interface (127.0.0.1 or localhost) and thus only you will be able to make requests to CouchDB, nobody else.
- If you don't like that, you can create specific admin users with a username and password as their credentials.
- You could also add **validation functions** to your design documents that will make sure that the right people make the right changes.
- **Securing CouchDB is outside the scope of this lecture... The same applies to web programming with CouchDB.**